

Networks (Tutorial). Week 4

Shinnazar Seytnazarov, PhD

Innopolis University

s.seytnazarov@innopolis.ru

January 28, 2022

Topic of the lecture

- Electronic Mail
- SMTP
- POP3
- IMAP4
- DNS
- P2P Applications
- Socket Programming with UDP and TCP

Topic of the tutorial

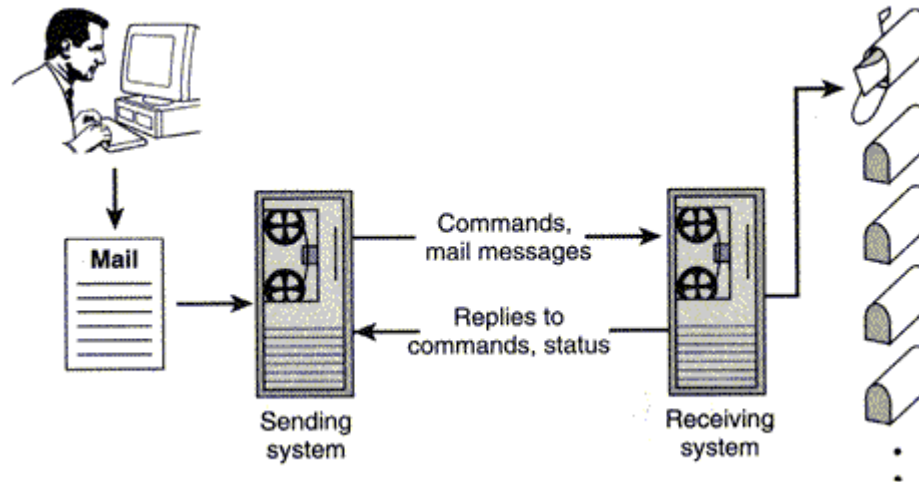
- Email Protocol Suit
- TCP Socket Programming

Topic of the lab

- TCP Socket Programming

Email

- E-mail has become an **essential means** to communicate.



The Mail Protocol Suite

- The most common mail protocols include:
 - SMTP (Simple Mail Transfer Protocol)
 - POP3 (Post Office Protocol)
 - IMAP4 (Internet Message Access Protocol)
 - MIME (Multipurpose Internet Mail Extensions)


Standards to Exchange Information

- ASCII stands for the "American Standard Code for Information Interchange".
- It was designed in the early 60's, as a standard character set for computers and electronic devices.
- ASCII is a 7-bit character set containing 128 characters.
- It contains the numbers from 0-9, the upper and lower case English letters from A to Z, and some special characters.
- The character sets used in modern computers and over the Internet.

ASCII Character Set

- For Example
 - ASCII complaint computer use:

65 to represent A



0100 0001

This standard is unable to support many languages characters. For instance: Russian

Source: <http://www.mathcs.emory.edu/~cheung/Courses/170/Syllabus/05/char1.html>

Unicode

- Unicode is a **generalization of ASCII** using **16-bits**.
- It allows up to **65536 distinct members**.
- Unicode **includes ISO 8859-1** as the **first 256 members** and then
 - Cyrillic characters
 - Both traditional and simplified Chinese,
 - Japanese and Thai characters,
 - Korean and many more languages
- **Few email tools** support Unicode today.

Rich Text Format (RTF)

- RTF is ASCII text with some additions to allow specification of formatting information:
 - Type face, font size and color.
- Sending email in RTF have a more professional appearance
- Two competing rich text standards
 - Microsoft's RTF (used in MS Word and Exchange)
 - MIME (Multipurpose Internet Mail Extension)
- **NOTE:** Good systems can keep track which system supports rich text format and automatically convert into simple ASCII for anyone else.

Multipurpose Internet Mail Extension (MIME)

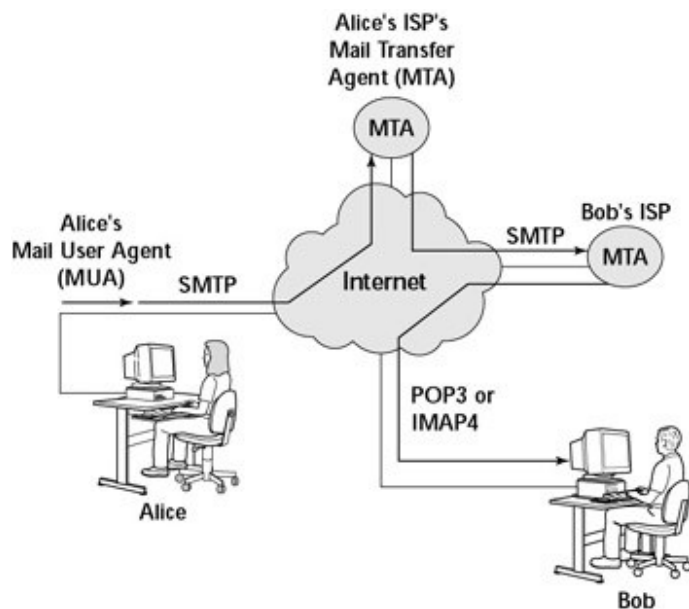
- MIME is an **Internet standard** that extends the format of email to support:
 - **ASCII** characters
 - **Binary:** Non-text – audio, video, images, application programs etc.
- MIME was designed **mainly for SMTP**
- The content types defined by MIME standards are also of importance in **communication protocols outside of email**, such as
 - **HTTP** for the World Wide Web.

Email Attachment

- An email attachment is a **binary file**
 - **Special Case:** It could be an **ASCII**
- It is **not in the main body part** of the message
- An **attachment tags along with the main body** part goes thorough the mail distribution system
- On the recipient's end, there is **an indication that attachments** have been received.

Message Transfer Agent (MTA)

- MTA is also known as [an email server](#)
- It is responsible for [collecting mail](#) from and [distributing mail](#) to user agents



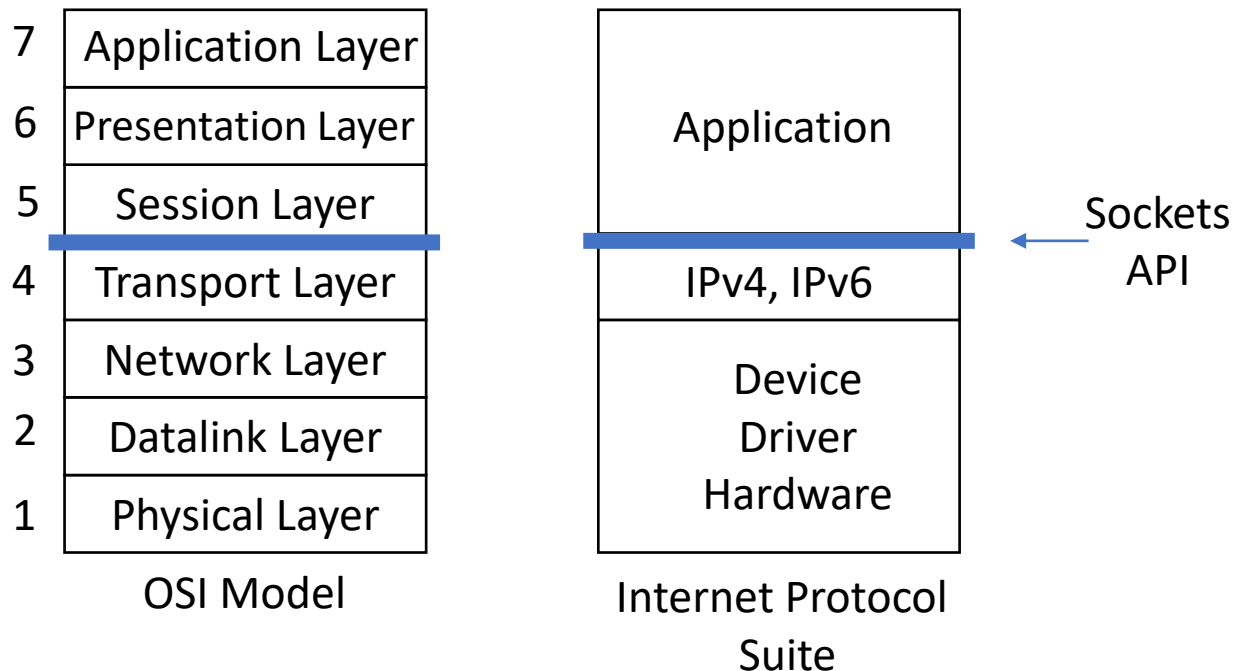
- The exact method used for [mail transport](#) depends on the underlying network.
- In [TCP/IP networks](#): the mail-transport [agent delivers mail](#) using SMTP.

Socket

- We can **communicate with application** which is located in **different systems** through communication method named **Socket**.
- Socket is the **end point of two way communication** link between two programs running on network.
- You can **also communicate between two processes** through socket on same machine but mostly it is used for **communicating over network**.
- Socket is supported by **different operating systems** like window, MAC, Linux, Unix etc.

Socket

- Socket **programming interface** in relation to the protocol stack



Types of Socket

- Three types of socket
 - Stream socket
 - Datagram socket
 - Raw socket

Types of Socket

- **Stream socket**

- Reliable
- Connection-oriented
- TCP

- **Datagram socket**

- Unreliable
- Connectionless
- UDP

- **RAW socket**

- It provides RAW data transfer directly over IP Protocol (no transport layer)

Stream Socket

- When a computer is connected to the internet, it gets an IP address.
- This address uniquely identifies a computer on the Internet.
- In order to connect one to another, it must know the IP address of the host machine.
- In addition to the IP address, the process must know what “port” the other process is listening on.

Stream Socket

- **TCP Connection:** direct virtual pipe between client's socket and server's connection socket.
- A server is the process that is constantly running, listening for socket connections.
- The client process, attempts to make contact with the server, when needed.

TCP client/server Model Description

1. Client **initiates** contact with server.

2. Server **must be ready**
 - Not dormant (running).
 - Have a socket that welcomes initial contact from client process
 - Welcoming socket

3. Client initiates **TCP connection** to the server
 - Creates client socket (address of server process, port number).


TCP client/server Model Description

4. Three-way **handshake**

- Client knocks server welcoming socket.
- Server hears the knock and creates new socket to particular client (connection socket).
- TCP connection exist between client's socket and server's new socket.

5. Client sends data via its client socket to server via **connection socket** and TCP guarantees **the delivery of data**

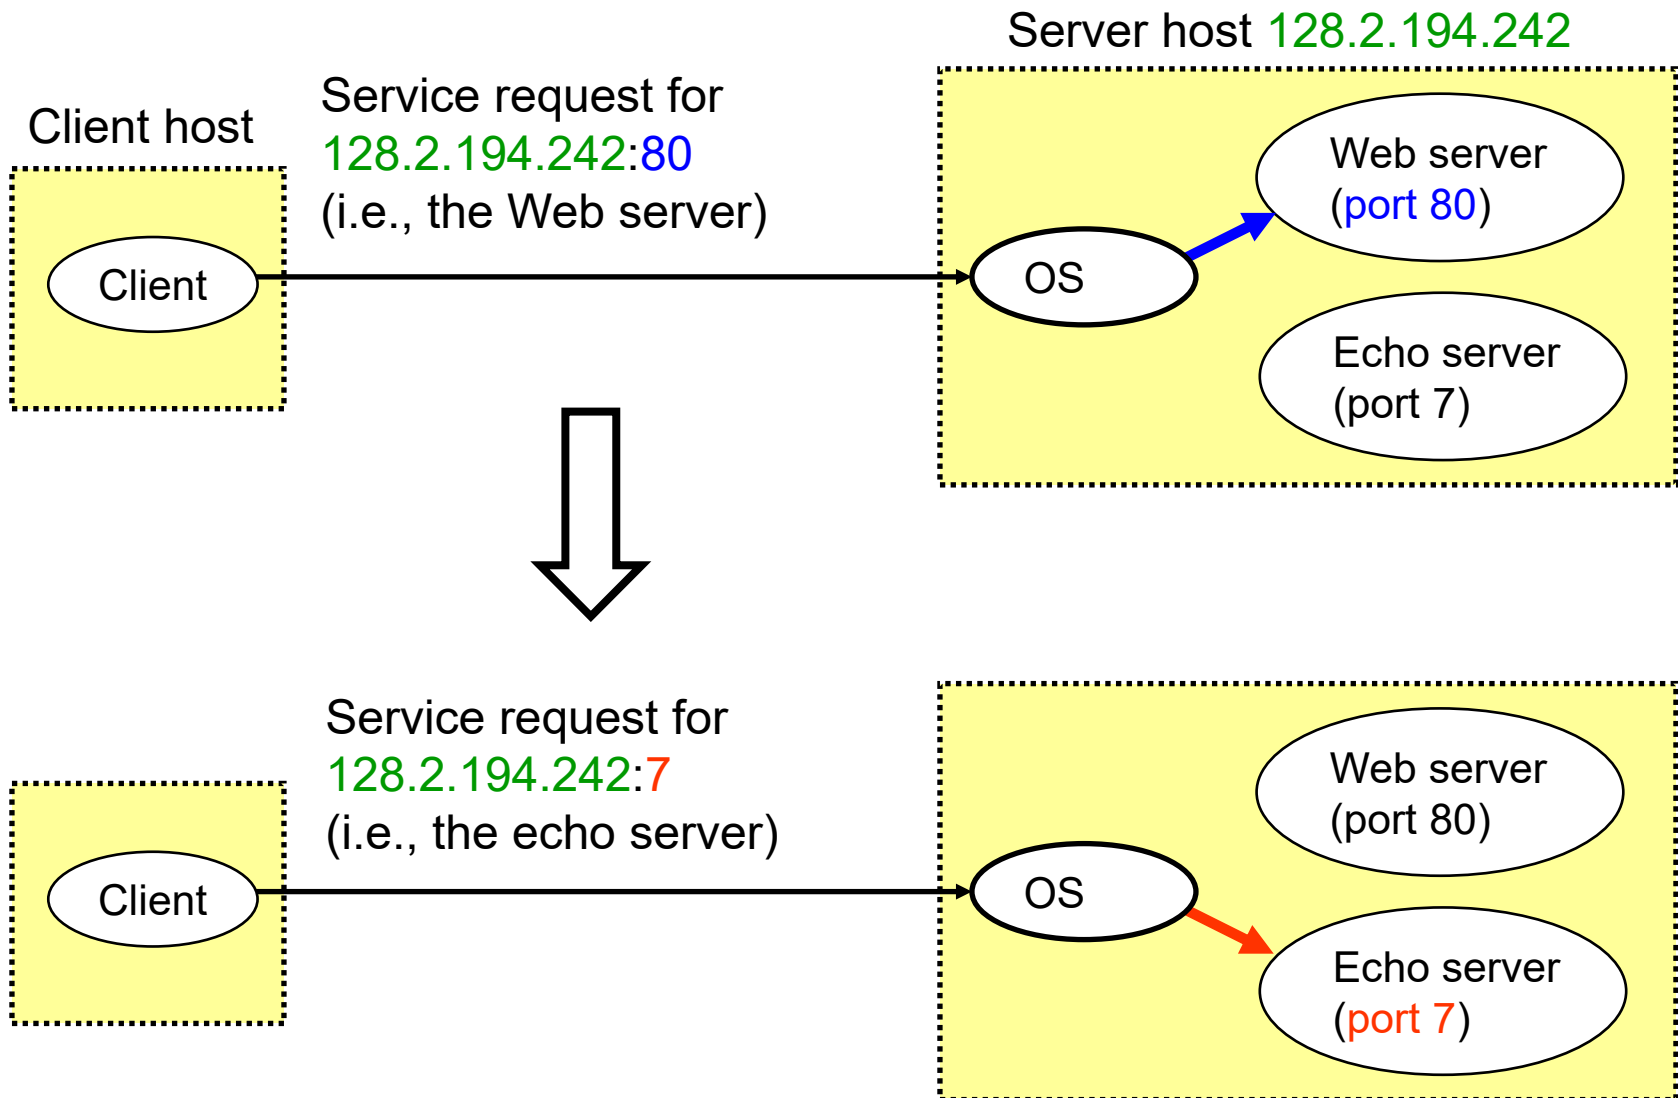
- It means that TCP provides reliable byte-stream service between client and service processes.



An Analogy!!

- **Socket:** telephone
- **Bind:** assign telephone number to a telephone
- **Listen:** turn on the ringer so that you can hear the phone call
- **Connect:** dial a phone number
- **Accept:** answer the phone
- **Read/write:** talking
- **Close:** hang up

Using Ports to Identify Services



UNIX Socket API

- In UNIX, everything is like a file
 - All input is like reading a file
 - All output is like writing a file
- API implemented as system calls
 - For Example: connect, send, recv, close, ...

Steps involved in establishing a socket on the client

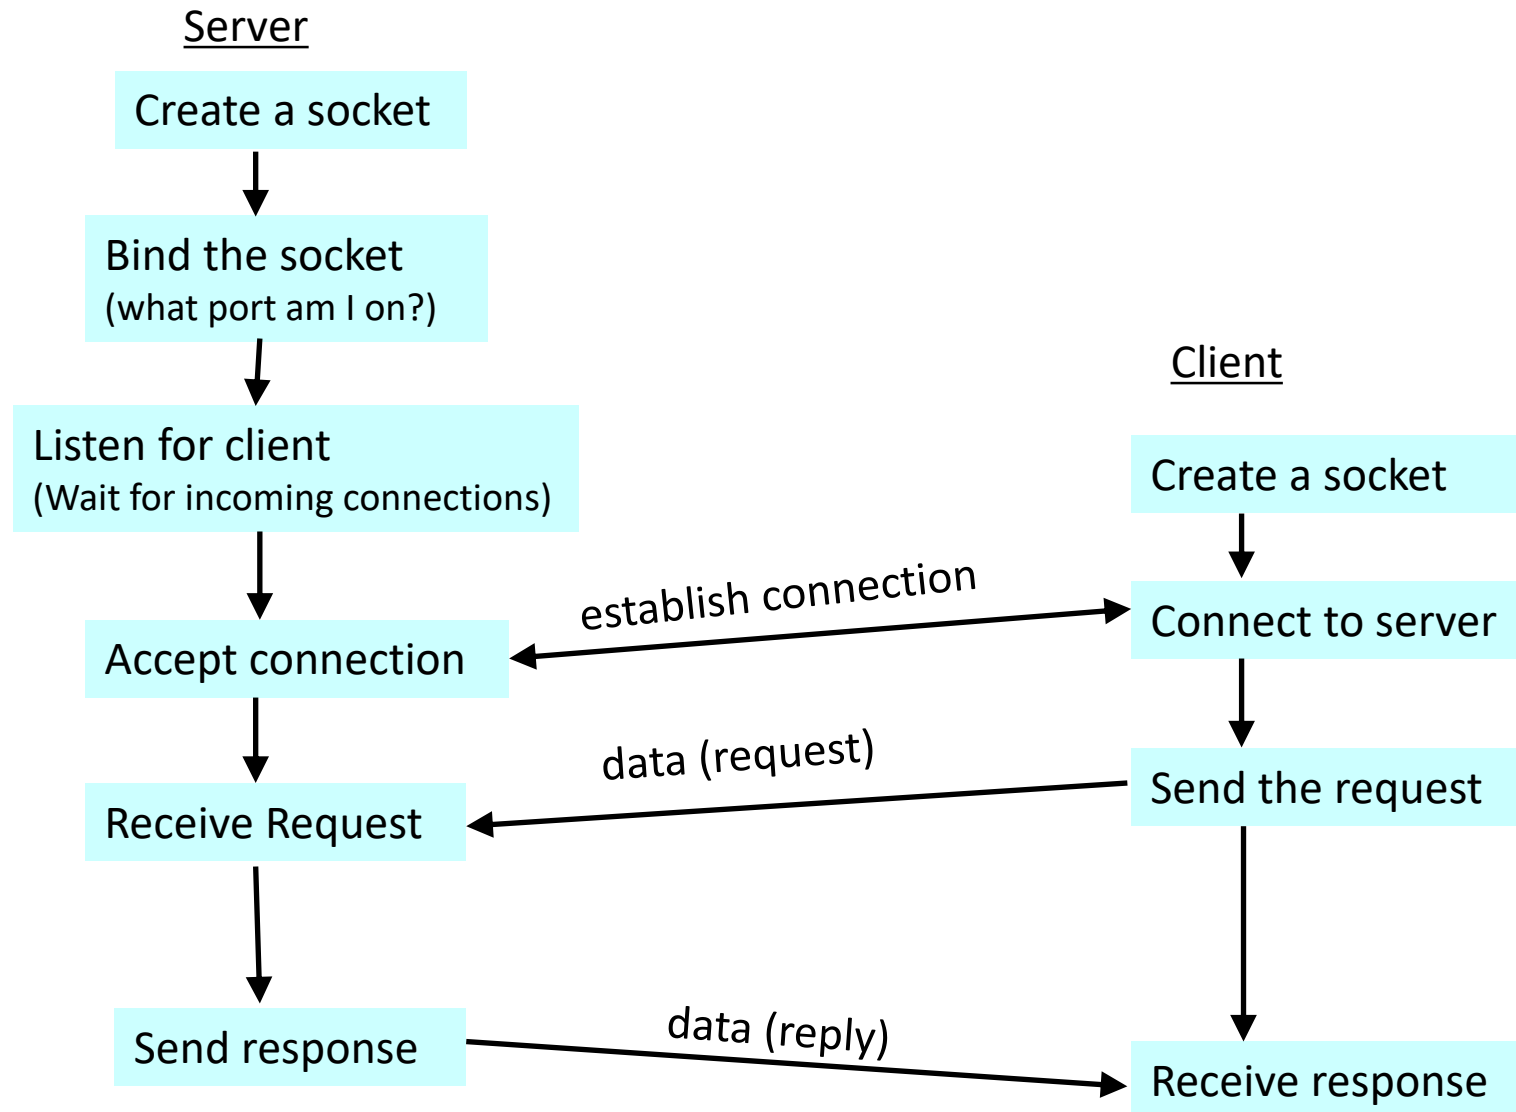
- Client side steps to create socket
 - Create a socket with the `socket()` system call
 - Connect the socket to the address of the server using the `connect()` system call
 - Send and receive data using `send()` and `recv()` system calls.

Steps involved in establishing a socket on the server

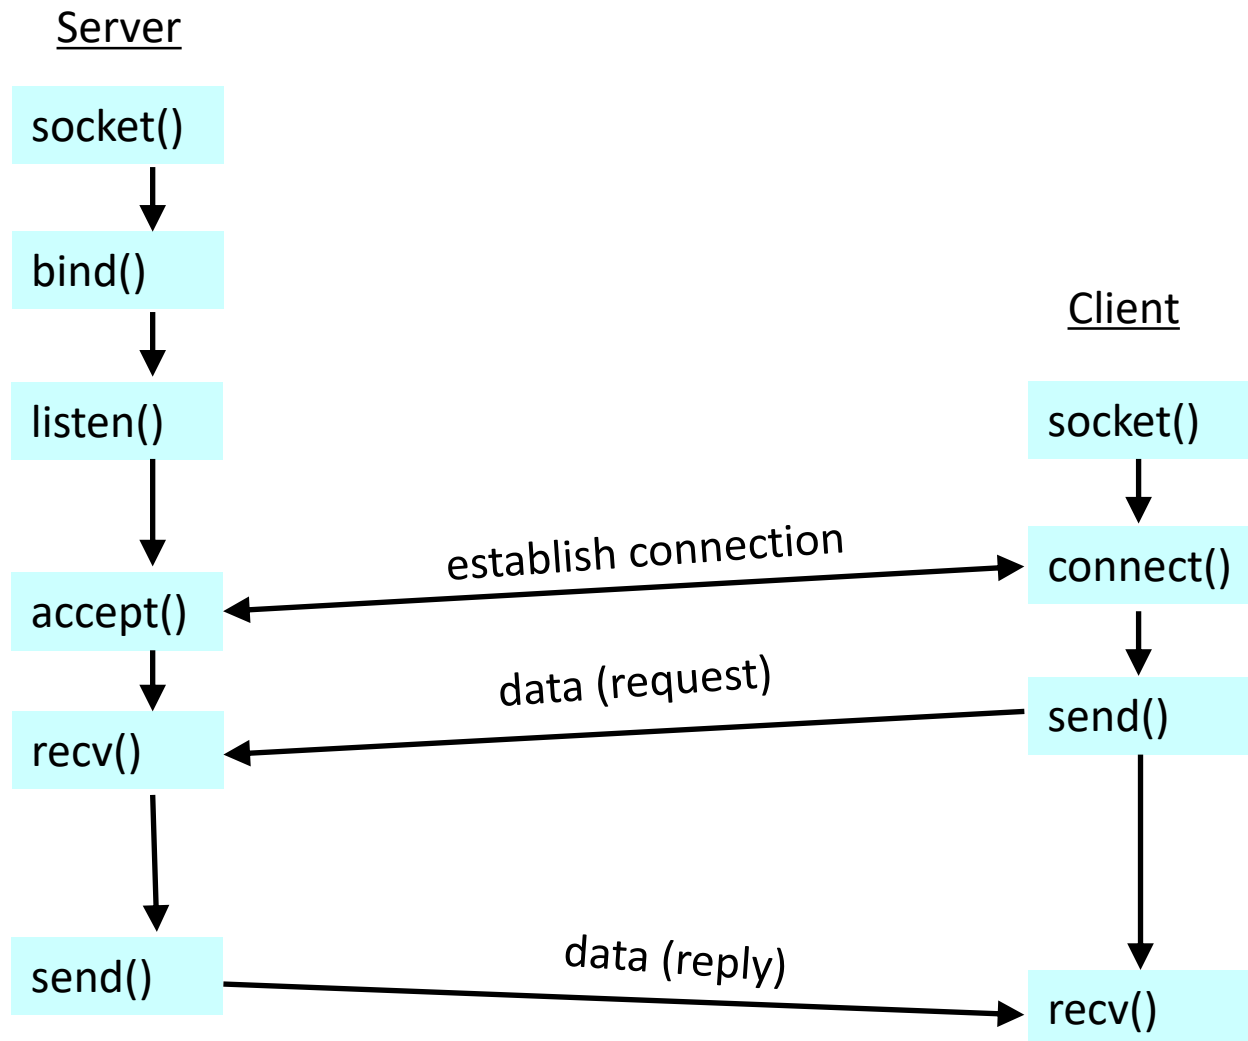
- Server side steps to create socket
 - Create a socket with the `socket()` system call
 - Bind the socket to an address using the `bind()` system call. For a server socket on the Internet, an address consists of a port number on the host machine.
 - Listen for connections with the `listen()` system call
 - Accept a connection with the `accept()` system call.
 - Send and receive data

Client-Server Communication

Stream Sockets (TCP): Connection-oriented



Connection-oriented Example (Stream Sockets -TCP)



Client: Learning Server Address/Port

- Server typically known by **name and service**
 - E.g., “www.cnn.com” and “http”
- Need to **translate into IP** address and **port #**
 - E.g., “64.236.16.20” and “80”
- Get address info with given host name and service
 - **int getaddrinfo(char *node,**
 char *service
 struct addrinfo *hints,
 struct addrinfo **result)
 - ***node**: host name (e.g., “www.cnn.com”) or IP address
 - ***service**: port number or service listed in */etc/services* (e.g. ftp)
 - **hints**: points to a *struct addrinfo* with known information

Client: Learning Server Address/Port (cont.)

- Data structure to host address information

```
struct addrinfo {
    int      ai_flags;
    int      ai_family;      //e.g. AF_INET for IPv4
    int      ai_socktype;    //e.g. SOCK_STREAM for TCP
    int      ai_protocol;    //e.g. IPPROTO_TCP
    size_t   ai_addrlen;
    char     *ai_canonname;
    struct sockaddr *ai_addr; // point to sockaddr struct
    struct addrinfo *ai_next;
}
```

- Example

```
hints.ai_family = AF_UNSPEC;      // don't care IPv4 or IPv6
hints.ai_socktype = SOCK_STREAM; // TCP stream sockets
int status = getaddrinfo("www.cnn.com", "80", &hints, &result);
// result now points to a linked list of 1 or more addrinfos
// etc.
```

Client: Creating a Socket

- Creating a socket
 - `int socket(int domain, int type, int protocol)`
 - Returns a file descriptor (or handle) for the socket
- Domain: Protocol family
 - PF_INET for IPv4
 - PF_INET6 for IPv6
- Type: Semantics of the communication
 - SOCK_STREAM: reliable byte stream (TCP)
 - SOCK_DGRAM: message-oriented service (UDP)
- Protocol: Specific protocol
 - UNSPEC: unspecified
 - (PF_INET and SOCK_STREAM already implies TCP)
- Example


```
sockfd = socket(result->ai_family, result->ai_socktype, result->ai_protocol);
```


Client: Connecting Socket to the Server

- Client contacts the server to establish connection
 - Associate the socket with the server address/port
 - Acquire a local port number (assigned by the OS)
 - Request connection to server, who hopefully accepts
 - connect is **blocking**

- Establishing the connection
 - `int connect(int sockfd,
 struct sockaddr *server_address,
 socketlen_t addrlen)`

 - **Args:** socket descriptor, server address, and address size
 - Returns 0 on success, and -1 if an error occurs

- For Example:


```
connect(sockfd, result->ai_addr, result->ai_addrlen);
```

Client: Sending Data

- **Sending data**

- `int send(int sockfd, void *msg, size_t len, int flags)`

- **Arguments**

- Socket descriptor, pointer to buffer of data to send, and length of the buffer

- **Returns**

- The number of bytes written, and -1 on error

- **Note**

- Send is blocking: return only after data is sent
 - Write short messages into a buffer and send once

Client: Receiving Data

- **Receiving data**

```
int recv(int sockfd, void *buf, size_t len, int flags)
```

- **Arguments**

- Socket descriptor, pointer to buffer to place the data, size of the buffer, and flags

- **Returns**

- The number of characters read (where 0 implies “end of file”), and -1 on error

Server: Server Preparing its Socket

- Server creates a socket **and binds address/port**
 - Server creates a socket, just like the client does
 - Server associates the socket with the port number
- Create a socket
 - `int socket(int domain, int type, int protocol)`
- Bind socket to the local address and port number
 - `int bind(int sockfd, struct sockaddr *my_addr, socklen_t addrlen)`

Server: Allowing Clients to Wait

- Many client **requests** may arrive
 - Server cannot handle them all at the same time
 - Server could reject the requests, or let them wait
- Define how many connections can be pending
 - **`int listen(int sockfd, int backlog)`**
 - **Arguments:** socket descriptor and acceptable backlog
 - **Returns** a 0 on success, and -1 on error
 - Listen is **non-blocking**: returns immediately
- **What if too many clients arrive?**
 - Some requests don't get through
 - The Internet makes no promises...
 - And the client can always **try again**



Server: Accepting Client Connection

- Now all the server can do is wait...
 - Waits for connection request to arrive
 - **Blocking** until the request arrives
 - And then accepting the new request



- Accept a new connection from a client
 - `int accept(int sockfd, struct sockaddr *addr, socketlen_t *addrlen)`
 - **Arguments:**
 - sockfd, structure that will provide client address and port, and length of the structure
 - **Returns**
 - Descriptor of socket for this new connection

Client and Server: Cleaning House

- Once the connection is open
 - Both sides can read and write
 - Two unidirectional streams of data
 - In practice, client writes first, and server reads
 - ... then server writes, and client reads, and so on
- Closing down the connection
 - Either side can close the connection
 - ... using the `int close(int sockfd)`
- What about the data still “in flight”
 - Data in flight still reaches the other end
 - So, server can `close()` before client finishes reading

Server: One Request at a Time?

- Serializing requests is **inefficient**
 - Server can process just one request at a time
 - All other clients must wait until previous one is done

- May **need to time share** the server machine
 - Alternate between servicing different requests
 - Do a **little work on one request**, then switch when you are waiting for some other resource
 - **“Nonblocking I/O”**

 - Or, use a **different process/thread** for each request
 - Allow OS to share the CPU(s) across processes

 - Or, some **hybrid of these two** approaches

Handle Multiple Clients using fork()

- Steps to handle multiple clients
 - Go to a loop and accept connections using `accept()`
 - After a connection is established, call `fork()`
 - To create a new child process to handle it
 - Go back to listen for another socket in the parent process
 - `close()` when you are done.

Real Clients and Servers?

- Apache Web server
 - Open source server first released in 1995
 - Software available online at <http://www.apache.org>
- Mozilla Web browser
 - <http://www.mozilla.org/developer/>
- Sendmail
 - <http://www.sendmail.org/>
- BIND Domain Name System
 - Client resolver and DNS server
 - <http://www.isc.org/index.pl?/sw/bind/>
- ...

Acknowledgements

- This tutorial was prepared by M.Fahim, G.Succi, and A.Tormasov

Reference

- This tutorial is based on the following documents as well as lecture materials over the internet.
 - <http://users.cs.cf.ac.uk/Dave.Marshall/Internet/node86.html>
 - <http://alumni.cs.ucr.edu/~ecegelal/TAw/socketTCP.pdf>