

Databases 2022

Darko Bozhinoski,
Ph.D. in Computer Science

Agenda

- Data Definition Language (DDL) (Recap)
- Data Manipulation Language (DML)
- SQL Advanced Concepts

Relational Algebra Operators

Basic Operators

Extended / Derived Operators

Two types of operations:

- set operations from set theory
- operations developed specifically for relational databases

Unary Operators

Binary Operators

Projection Operator (π)

Selection Operator (σ)

Rename Operator (ρ)

Union Operator (\cup)

Cross Product Operator (\times)

Minus / Set Difference Operator ($-$)

Join Operator (\bowtie)

Division Operator ($/$ or \div)

Intersection Operator (\cap)

SQL ENVIRONMENT

- **Catalog**

- A set of schemas that constitute the description of a database

- **Schema**

- The structure that contains descriptions of objects created by a user (base tables, views, constraints)

- **Data Definition Language (DDL)**

- Commands that define a database, including creating, altering, and dropping tables and establishing constraints (CREATE, ALTER, DROP)

- **Data Manipulation Language (DML)**

- Commands that maintain and query a database (INSERT, UPDATE, DELETE, SELECT)

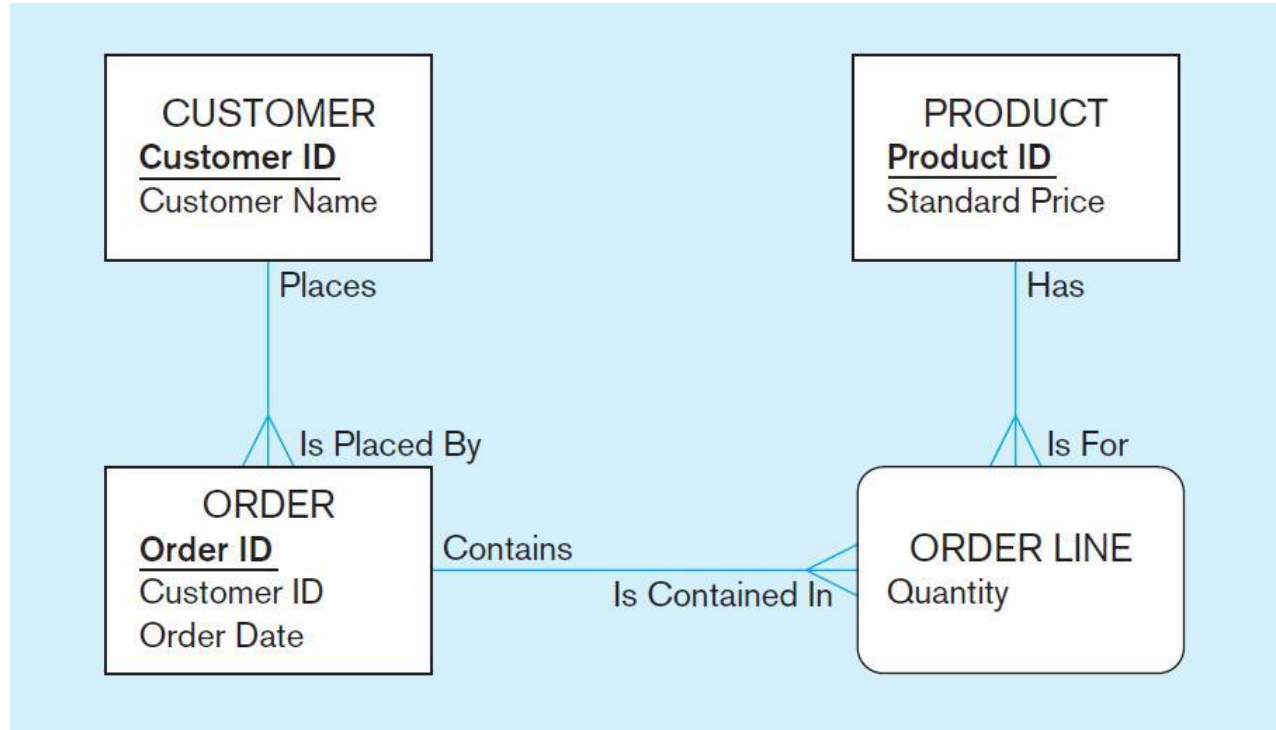
- **Data Control Language (DCL)**

- Commands that control a database, including administering privileges and committing data (GRANT, REVOKE)

STEPS IN TABLE CREATION

- ❖ Identify data types for attributes
- ❖ Identify columns that can and cannot be null
- ❖ Identify columns that must be unique (candidate keys)
- ❖ Identify primary key–foreign key mates
- ❖ Determine default values
- ❖ Identify constraints on columns (domain specifications)
- ❖ Create the table and associated indexes

Example DATA MODEL



SQL database definition commands

Overall table definitions

```
CREATE TABLE Customer_T
    (CustomerID          NUMBER(11,0)    NOT NULL,
     CustomerName        VARCHAR2(25)    NOT NULL,
     CustomerAddress     VARCHAR2(30),
     CustomerCity        VARCHAR2(20),
     CustomerState       CHAR(2),
     CustomerPostalCode  VARCHAR2(9),
    CONSTRAINT Customer_PK PRIMARY KEY (CustomerID));
```

```
CREATE TABLE Order_T
    (OrderID             NUMBER(11,0)    NOT NULL,
     OrderDate           DATE DEFAULT SYSDATE,
     CustomerID          NUMBER(11,0),
    CONSTRAINT Order_PK PRIMARY KEY (OrderID),
    CONSTRAINT Order_FK FOREIGN KEY (CustomerID) REFERENCES Customer_T(CustomerID));
```

```
CREATE TABLE Product_T
    (ProductID          NUMBER(11,0)    NOT NULL,
     ProductDescription  VARCHAR2(50),
     ProductFinish      VARCHAR2(20)
                        CHECK (ProductFinish IN ('Cherry', 'Natural Ash', 'White Ash',
                                                  'Red Oak', 'Natural Oak', 'Walnut')),
     ProductStandardPrice DECIMAL(6,2),
     ProductLineID      INTEGER,
    CONSTRAINT Product_PK PRIMARY KEY (ProductID));
```

```
CREATE TABLE OrderLine_T
    (OrderID            NUMBER(11,0)    NOT NULL,
     ProductID          INTEGER         NOT NULL,
     OrderedQuantity    NUMBER(11,0),
    CONSTRAINT OrderLine_PK PRIMARY KEY (OrderID, ProductID),
    CONSTRAINT OrderLine_FK1 FOREIGN KEY (OrderID) REFERENCES Order_T(OrderID),
    CONSTRAINT OrderLine_FK2 FOREIGN KEY (ProductID) REFERENCES Product_T(ProductID));
```

Controlling the values in attributes

```
CREATE TABLE Order_T
    (OrderID                NUMBER(11,0)    NOT NULL,
     OrderDate              DATE DEFAULT SYSDATE,
     CustomerID             NUMBER(11,0),
 CONSTRAINT Order_PK PRIMARY KEY (OrderID),
 CONSTRAINT Order_FK FOREIGN KEY (CustomerID) REFERENCES Customer_T(CustomerID));
```

Default value

```
CREATE TABLE Product_T
    (ProductID              NUMBER(11,0)    NOT NULL,
     ProductDescription      VARCHAR2(50),
     ProductFinish           VARCHAR2(20)
 CHECK (ProductFinish IN ('Cherry', 'Natural Ash', 'White Ash',
                          'Red Oak', 'Natural Oak', 'Walnut')),
     ProductStandardPrice   DECIMAL(6,2),
     ProductLineID          INTEGER,
 CONSTRAINT Product_PK PRIMARY KEY (ProductID));
```

Domain constraint

DATA INTEGRITY CONTROLS

- ❖ Referential integrity—constraint that ensures that foreign key values of a table must match primary key values of a related table in 1:M relationships
- ❖ Restricting:
 - Deletes of primary records
 - Updates of primary records
 - Inserts of dependent records

Key and Referential Integrity Constraints

The schema designer can specify an alternative action to be taken by attaching a referential triggered action clause to any foreign key constraint. The options include: SET NULL, CASCADE, and SET DEFAULT.

An option must be qualified with either ON DELETE or ON UPDATE.

```
CREATE TABLE EMPLOYEE
(
    ...,
    Dno INT NOT NULL DEFAULT 1,
    CONSTRAINT EMPPK
    PRIMARY KEY (Ssn),
    CONSTRAINT EMPSUPERFK
    FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn)
    ON DELETE SET NULL ON UPDATE CASCADE
    CONSTRAINT EMPDEPTFK
    FOREIGN KEY (Dno) REFERENCES DEPARTMENT(Dnumber)
    ON DELETE SET DEFAULT ON UPDATE CASCADE
);

CREATE TABLE DEPARTMENT
(
    ...,
    Mgr_ssn CHAR(9) NOT NULL DEFAULT '888665555',
    ...,
    CONSTRAINT DEPTPK
    PRIMARY KEY (Dnumber),
    CONSTRAINT DEPTSK
    UNIQUE (Dname),
    CONSTRAINT DEPTMGRFK
    FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn)
    ON DELETE SET DEFAULT ON UPDATE CASCADE
);

CREATE TABLE DEPT_LOCATIONS
(
    ...,
    PRIMARY KEY (Dnumber, Dlocation),
    FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber)
    ON DELETE CASCADE ON UPDATE CASCADE
);
```

SQL Query

Basic SQL Query

SELECT	<attributes>
FROM	<one or more relations>
WHERE	<conditions>

SQL Query to Relational Algebra Basic

```
SELECT Select-list  
FROM  $R_1, \dots, R_2 T_2, \dots$   
WHERE Where-condition
```

When the statement does not use subqueries in its where-condition, we can easily translate it into the relational algebra as follows:


$$\pi_{\text{Select-list}} \sigma_{\text{Where-condition}} (R_1 \times \dots \times \rho_{T_2}(R_2) \times \dots).$$

Simple SQL Query

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT *  
FROM Product  
WHERE category='Gadgets'
```



“selection”


PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks

Simple SQL Query

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT PName, Price,  
Manufacturer  
FROM Product  
WHERE Price > 100
```



PName	Price	Manufacturer
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi


“selection” and
“projection”

Notation

Input Schema

Product(PName, Price, Category, Manufacturer)

```
SELECT PName, Price, Manufacturer  
FROM   Product  
WHERE  Price > 100
```



Output Schema

Answer(PName, Price, Manufacturer)


The **LIKE** operator

```
SELECT  *  
FROM    Products  
WHERE   PName LIKE '%gizmo%'
```

- **s LIKE p**: pattern matching on strings
- **p** may contain two special symbols:
 - **%** = any sequence of characters
 - **_** = any single character

Eliminating Duplicates


```
SELECT DISTINCT category  
FROM   Product
```



Category
Gadgets
Photography
Household

Compare to:

```
SELECT category  
FROM   Product
```



Category
Gadgets
Gadgets
Photography
Household

Ordering the Results


```
SELECT  pname, price, manufacturer  
FROM    Product  
WHERE   category='gizmo' AND price > 50  
ORDER BY price, pname
```

Ties are broken by the second attribute on the ORDER BY list, etc.

Ordering is ascending, unless you specify the DESC keyword.

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi


```
SELECT DISTINCT category
FROM Product
ORDER BY category
```



?


PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT DISTINCT category  
FROM Product  
ORDER BY category
```



?


```
SELECT Category  
FROM Product  
ORDER BY PName
```



?


PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT DISTINCT category
FROM Product
ORDER BY category
```




?

```
SELECT Category
FROM Product
ORDER BY PName
```



?


```
SELECT DISTINCT category
FROM Product
ORDER BY PName
```




?

JOINS


LEFT JOIN



RIGHT JOIN



INNER JOIN



FULL JOIN




Table 1

Column 1	Column2
A	X
B	Y
C	Z
D	W

Table 2

Column2	Column3
X	XName
Y	YName
P	PName
O	OName

Difference between Cross Product and Full Outer Join?

Cartesian Product

Select * from Table1, Table2

Column 1	Column2	Column2	Column3
A	X	X	XName
A	X	Y	YName
A	X	P	PName
A	X	O	OName
B	Y	X	XName
B	Y	Y	YName
B	Y	P	PName
B	Y	O	OName
C	Z	X	XName
C	Z	Y	YName
C	Z	P	PName
C	Z	O	OName
D	W	X	XName
D	W	Y	YName
D	W	P	PName
D	W	O	OName

A cross join produces a cartesian product between the two tables, returning all possible combinations of all rows.

FULL OUTER JOIN

Column 1	Column2	Column2	Column3
A	X	X	XName
B	Y	Y	YName
C	Z	NULL	NULL
D	W	NULL	NULL
NULL	NULL	P	PName
NULL	NULL	O	OName

FULL OUTER JOIN:

- It includes all the rows from both the tables.
- Assigns NULL for unmatched fields.
- A combination of both left and right outer joins.

JOINS

Product (pname, price, category, manufacturer)

Company (cname, stockPrice, country)

Find all products under \$200 manufactured in Japan;
return their names and prices.

```
SELECT  PName, Price
FROM    Product, Company
WHERE   Manufacturer=CName AND Country='Japan'
        AND Price <= 200
```

JOINS (2)


Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

Company

Cname	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

```
SELECT PName, Price
FROM Product, Company
WHERE Manufacturer=CName AND Country='Japan'
AND Price <= 200
```



PName	Price
SingleTouch	\$149.99

JOINS (3)

Product (pname, price, category, manufacturer)

Company (cname, stockPrice, country)

Find all Russian companies that manufacture products both in the 'electronic' and 'toy' categories

```
SELECT  cname
```

```
FROM
```

```
WHERE
```

JOINS (4)

Product (pname, price, category, manufacturer)

Company (cname, stockPrice, country)

Find all countries that manufacture some product in the 'Gadgets' category.

```
SELECT  Country
FROM    Product, Company
WHERE   Manufacturer=CName AND Category='Gadgets'
```

JOINS (5)

Product


<u>Name</u>	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

Company

<u>Cname</u>	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

```
SELECT Country
FROM Product, Company
WHERE Manufacturer=CName AND Category='Gadgets'
```

What is
the problem ?
What's the
solution ?



Country
??
??

OUTER JOINS

Explicit joins in SQL = “inner joins”:

Product(name, category)

Purchase(prodName, store)

```
SELECT Product.name, Purchase.store
FROM    Product JOIN Purchase ON
        Product.name = Purchase.prodName
```

Same as:

```
SELECT Product.name, Purchase.store
FROM    Product, Purchase
WHERE   Product.name = Purchase.prodName
```

But Products that never sold will be lost !

OUTER JOINS

Left outer joins in SQL:

Product(name, category)

Purchase(prodName, store)

```
SELECT Product.name, Purchase.store  
FROM   Product LEFT OUTER JOIN Purchase ON  
        Product.name = Purchase.prodName
```

```
SELECT Product.name, Purchase.store
FROM   Product LEFT OUTER JOIN Purchase ON
       Product.name = Purchase.prodName
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Name	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz
OneClick	NULL

Example

Compute, for each product, the total number of sales in ‘September’

Product(name, category)

Purchase(prodName, month, store)

```
SELECT Product.name, count(*)  
FROM   Product, Purchase  
WHERE  Product.name =  
        Purchase.prodName  
        and Purchase.month = ‘September’  
GROUP BY Product.name
```

What’s wrong ?

Example

Compute, for each product, the total number of sales in ‘September’

Product(name, category)

Purchase(prodName, month, store)

```
SELECT Product.name, count(*)  
FROM    Product LEFT OUTER JOIN Purchase ON  
        Product.name = Purchase.prodName  
        and Purchase.month = ‘September’  
GROUP BY Product.name
```

Now we also get the products who sold in 0 quantity

Tuple Variables

Person(pname, address, worksfor)

Company(cname, address)

```
SELECT  DISTINCT pname, address  
FROM    Person, Company  
WHERE   worksfor = cname
```


Which address is this?

Tuple Variables

Person(pname, address, worksfor)

Company(cname, address)

```
SELECT  DISTINCT pname, address  
FROM    Person, Company  
WHERE   worksfor = cname
```




```
SELECT  DISTINCT Person.pname, Company.address  
FROM    Person, Company  
WHERE   Person.worksfor = Company.cname
```

```
SELECT DISTINCT R.A  
FROM R, S, T  
WHERE R.A=S.A OR R.A=T.A
```

What does it compute ?


```
SELECT DISTINCT R.A  
FROM R, S, T  
WHERE R.A=S.A OR R.A=T.A
```

What does it compute ?



Computes $R \cap (S \cup T)$

Subqueries Returning Relations

Company(name, city)

Product(pname, maker)

Purchase(id, product, buyer)

Return cities where one can find
companies that manufacture
products bought by Ivan Ivanov

```
SELECT Company.city
FROM Company
WHERE Company.name IN
      (SELECT Product.maker
       FROM Purchase , Product
       WHERE Product.pname=Purchase.product
        AND Purchase .buyer = 'Ivan Ivanov');
```

Subqueries Returning Relations

```
SELECT Company.city  
FROM    Company, Product, Purchase  
WHERE   Company.name= Product.maker  
          AND Product.pname = Purchase.product  
          AND Purchase.buyer = 'Ivan Ivanov'
```

Subqueries Returning Relations

```
SELECT Company.city  
FROM    Company, Product, Purchase  
WHERE   Company.name= Product.maker  
          AND Product.pname = Purchase.product  
          AND Purchase.buyer = 'Ivan Ivanov'
```

Beware of duplicates !

Removing Duplicates

```
SELECT DISTINCT Company.city
FROM   Company
WHERE  Company.name IN
      (SELECT Product.maker
       FROM   Purchase , Product
       WHERE  Product.pname=Purchase.product
              AND Purchase .buyer = 'Joe Blow');
```

```
SELECT DISTINCT Company.city
FROM   Company, Product, Purchase
WHERE  Company.name= Product.maker
      AND Product.pname = Purchase.product
      AND Purchase.buyer = 'Joe Blow'
```

Subqueries Returning Relations

You can also use: $s > \text{ALL } R$

$s > \text{ANY } R$

EXISTS R

Product (pname, price, category, maker)

Find products that are more expensive than all those produced
By “Gizmo-Works”

```
SELECT name
```

```
FROM Product
```

```
WHERE price > ALL (SELECT price
```

```
FROM Purchase
```

```
WHERE maker='Gizmo-Works')
```

Question

- Can we express this query as a single SELECT-FROM-WHERE query, without subqueries ?

Reading Material

- Fundamentals of Database Systems. Ramez Elmasri and Shamkant B. Navathe. Pearson. **Chapter 6. and Chapter 7.**
- SQL Tutorial: <https://www.w3schools.com/sql/default.asp>

Q & A

Three light-colored wooden blocks are arranged in a row on a dark wooden surface. The first block has a black letter 'Q', the second block has a black ampersand '&', and the third block has a black letter 'A'. The background is a soft-focus green and yellow bokeh.