



# Networks Lecture 3

Paolo Ciancarini

Innopolis University

January 24, 2022

# Source of the material

- This lecture is based on the following two resources
  - Chapter 15 of Computer Science Illuminated (3rd Edition) by Nell Dale and John Lewis
  - Chapter 1, 2 of Computer Networking: A Top-Down Approach (8th edition) by Jim Kurose and Keith Ross
  - The material is aligned and add/deleted according to the need of the students.

# Topic of the lecture

- Principles of Network Applications
- The transport service
- The Web and HTTP
- Cookies
- Web caching

# Addressing Processes

- To receive messages, process must have *identifier*
- Host device has unique 32-bit IP address
- **Q:** does IP address of host on which process runs suffice for identifying the process?
- **A:** no, many processes can be running on same host
- *Identifier* includes both **IP address** and **port numbers** associated with process on host.
- Example port numbers:
  - HTTP server: 80
  - Mail server: 25
- To send HTTP message to web server:
  - **IP address:** 128.119.245.12
  - **Port number:** 80

# App-layer protocol defines

- Types of messages exchanged
  - e.g., request, response
- Message syntax
  - what fields in messages & how fields are delineated
- Message semantics
  - meaning of information in fields
- Rules for when and how processes send & respond to messages

## Open protocols

defined in RFCs

allows for interoperability

e.g., HTTP, SMTP

## Proprietary protocols

e.g., Skype

# What transport service does an app need?

- **Data integrity**

- Some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- Other apps (e.g., audio) can tolerate some loss

- **Timing**

- Some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

- **Throughput**

- Some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- Other apps (“elastic apps”) make use of whatever throughput they get

- **Security**

- Encryption, data integrity, ...

## Requirements of selected apps: Common Apps

| application           | data loss     | throughput                                | time sensitive  |
|-----------------------|---------------|---|-----------------|
| file transfer         | no loss       | elastic                                   | no              |
| e-mail                | no loss       | elastic                                   | no              |
| Web documents         | no loss       | elastic                                   | no              |
| real-time audio/video | loss-tolerant | audio: 5kbps-1Mbps<br>video: 10kbps-5Mbps | yes, 100's msec |
| stored audio/video    | loss-tolerant | same as above                             | yes, few secs   |
| interactive games     | loss-tolerant | few kbps up                               | yes, 100's msec |
| text messaging        | no loss       | elastic                                   | yes and no      |

# Services provided by transport protocol to apps

## ***TCP Service:***

- *Reliable transport* between sending and receiving processes
- *Flow control*: the sender will not overwhelm the receiver
- *Congestion control*: throttle sender when network overloaded
- *Does not provide*: timing, minimum throughput guarantee, security
- *Connection-oriented*: setup required between client and server processes

## ***UDP service:***

- *Fast but unreliable data transfer* between sending and receiving process
- *Does not provide*: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup



# Internet apps: Application, Transport Protocols

| <b>application</b>     | <b>application layer protocol</b>                          | <b>underlying transport protocol</b> |
|------------------------|--|--------------------------------------|
| e-mail                 | SMTP [RFC 5321]  | TCP                                  |
| remote terminal access | Telnet [RFC 854]   | TCP                                  |
| Web                    | HTTP [RFC 7230]  | TCP                                  |
| file transfer          | FTP [RFC 959]  | TCP                                  |
| streaming multimedia   | HTTP (e.g., YouTube),<br>DASH                              | TCP or UDP                           |
| Internet telephony     | SIP[RFC 3261],<br>RTP[RFC 3550],<br>proprietary (eg Skype) | TCP or UDP                           |

# TCP vs UDP

## TCP



- Slower but more reliable transfers
- Typical Applications:
  - File Transfer Protocol (FTP)
  - Web Browsing
  - Email



**unicast**

## UDP



- Faster but not guaranteed transfers ("best effort")
- Typical Applications:
  - Live Streaming
  - Online Games
  - VoIP



**unicast**



**multicast**



**broadcast**

# Securing TCP

- TCP & UDP
  - No encryption
  - Cleartext passwords sent into socket traverse Internet in cleartext
- SSL (Secure Socket Layer)
  - Provides encrypted TCP connection
  - Data integrity
  - End-point authentication
- SSL is at app layer
  - Apps use SSL libraries, which “talk” to TCP
- SSL socket API
  - Clear text passwords sent into socket traverse Internet encrypted
  - Remark: Now SSL is called TLS Transport Layer Security
  - HTTP Over TLS is RFC 2818

# The Application Layer

- Principles of network applications
- Web and HTTP
- FTP

# Web and HTTP

*First, a review...*

- *A web page* includes some *objects* and some *links*
- An object can be a HTML file, JPEG image, Java applet, audio file,...
- A web page consists of *base HTML-file* which includes *several referenced objects*
- each object is addressable by a *URL*, e.g.,

`www.someschool.edu/someDept/pic.gif`

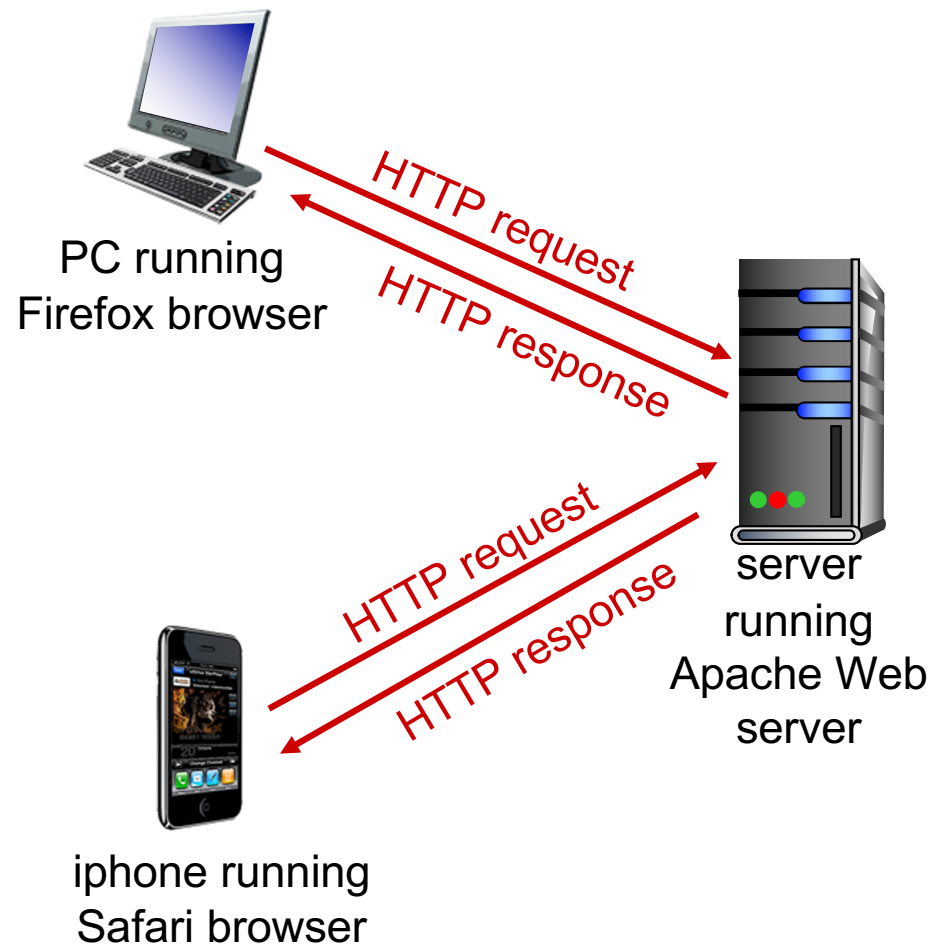
host name

path name

# HTTP Overview – 1/2

## HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
  - *client*: browser that requests, receives, (using HTTP protocol) and “displays” Web objects
  - *server*: Web server sends (using HTTP protocol) objects in response to requests



# HTTP Overview – 2/2

- *Uses TCP:*
  - Client initiates TCP connection (creates socket) to server, port 80
  - Server accepts TCP connection from client
  - HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
  - TCP connection closed
- *HTTP is “stateless”*
  - Server maintains no information about past client requests
- *Protocols that maintain “state” are complex!*
  - Past history (state) must be maintained
  - If server/client crashes, their views of “state” may be inconsistent, must be reconciled

# HTTP Connections

## *Non-persistent HTTP*

- At most one object sent over TCP connection
- Connection then closed
- Downloading multiple objects required multiple connections

## *Persistent HTTP*

- Multiple objects can be sent over a single TCP connection between the client and the server



# Non-persistent HTTP

Suppose user enters URL:

**`www.someSchool.edu/someDepartment/home.index`**

(contains text, references  
to 10 jpeg images)

- 1a. HTTP client initiates TCP  
connection to HTTP server  
(process) at `www.someSchool.edu`  
on port 80

time

# Non-persistent HTTP

Suppose user enters URL:

**`www.someSchool.edu/someDepartment/home.index`**

(contains text, references  
to 10 jpeg images)

1a. HTTP client initiates TCP  
connection to HTTP server  
(process) at `www.someSchool.edu`  
on port 80

1b. HTTP server at host  
`www.someSchool.edu` waiting  
for TCP connection at port 80.  
“accepts” connection, notifying  
client

time  
↓

# Non-persistent HTTP

Suppose user enters URL:

**`www.someSchool.edu/someDepartment/home.index`**

(contains text, references  
to 10 jpeg images)

1a. HTTP client initiates TCP  
connection to HTTP server  
(process) at `www.someSchool.edu`  
on port 80

1b. HTTP server at host  
`www.someSchool.edu` waiting  
for TCP connection at port 80.  
“accepts” connection, notifying  
client

2. HTTP client sends HTTP  
*request message* (containing  
URL) into TCP connection  
socket. Message indicates that  
client wants object  
`someDepartment/home.index`

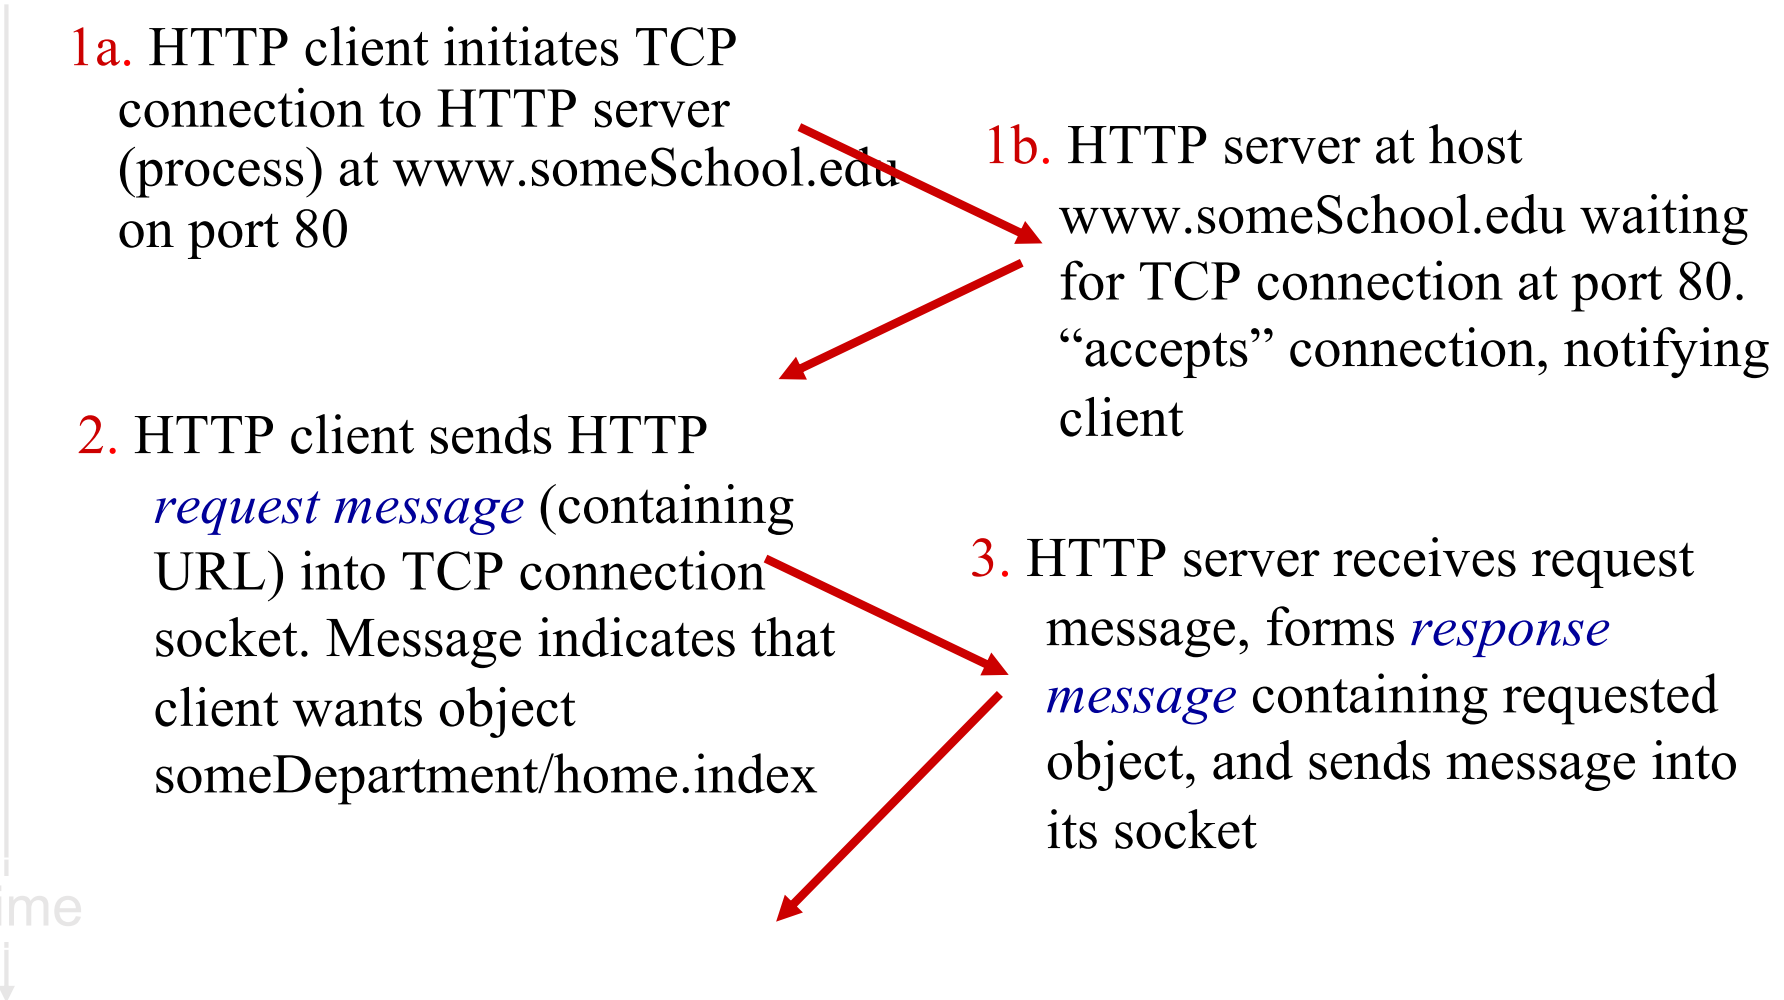
time  
↓

# Non-persistent HTTP

Suppose user enters URL:

**`www.someSchool.edu/someDepartment/home.index`**

(contains text, references  
to 10 jpeg images)

- 
- ```
sequenceDiagram
    participant Client
    participant Server
    Note over Client: 1a. HTTP client initiates TCP connection to HTTP server (process) at www.someSchool.edu on port 80
    Note over Server: 1b. HTTP server at host www.someSchool.edu waiting for TCP connection at port 80. "accepts" connection, notifying client
    Client->>Server: 2. HTTP client sends HTTP request message (containing URL) into TCP connection socket. Message indicates that client wants object someDepartment/home.index
    Server->>Client: 3. HTTP server receives request message, forms response message containing requested object, and sends message into its socket
```
- 1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80
- 1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. “accepts” connection, notifying client
2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`
3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

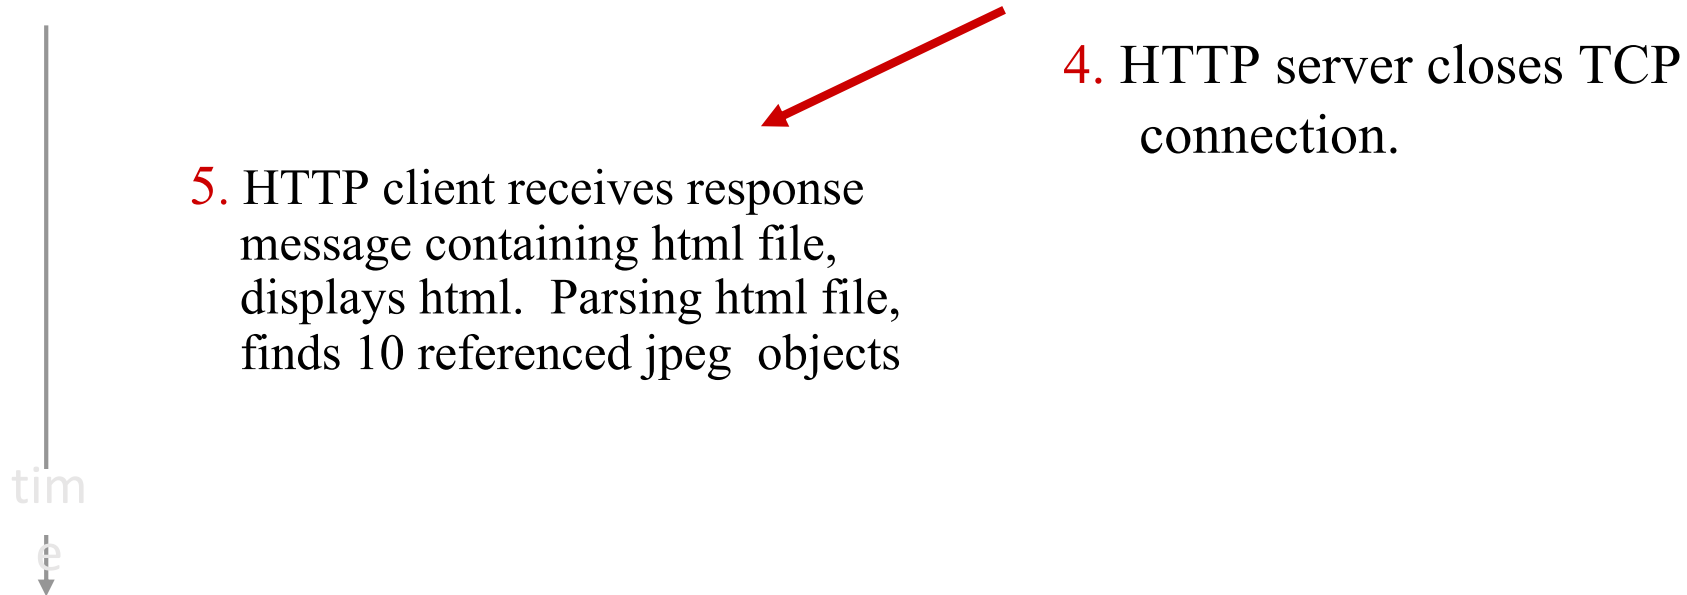
time  
↓

# Non-persistent HTTP

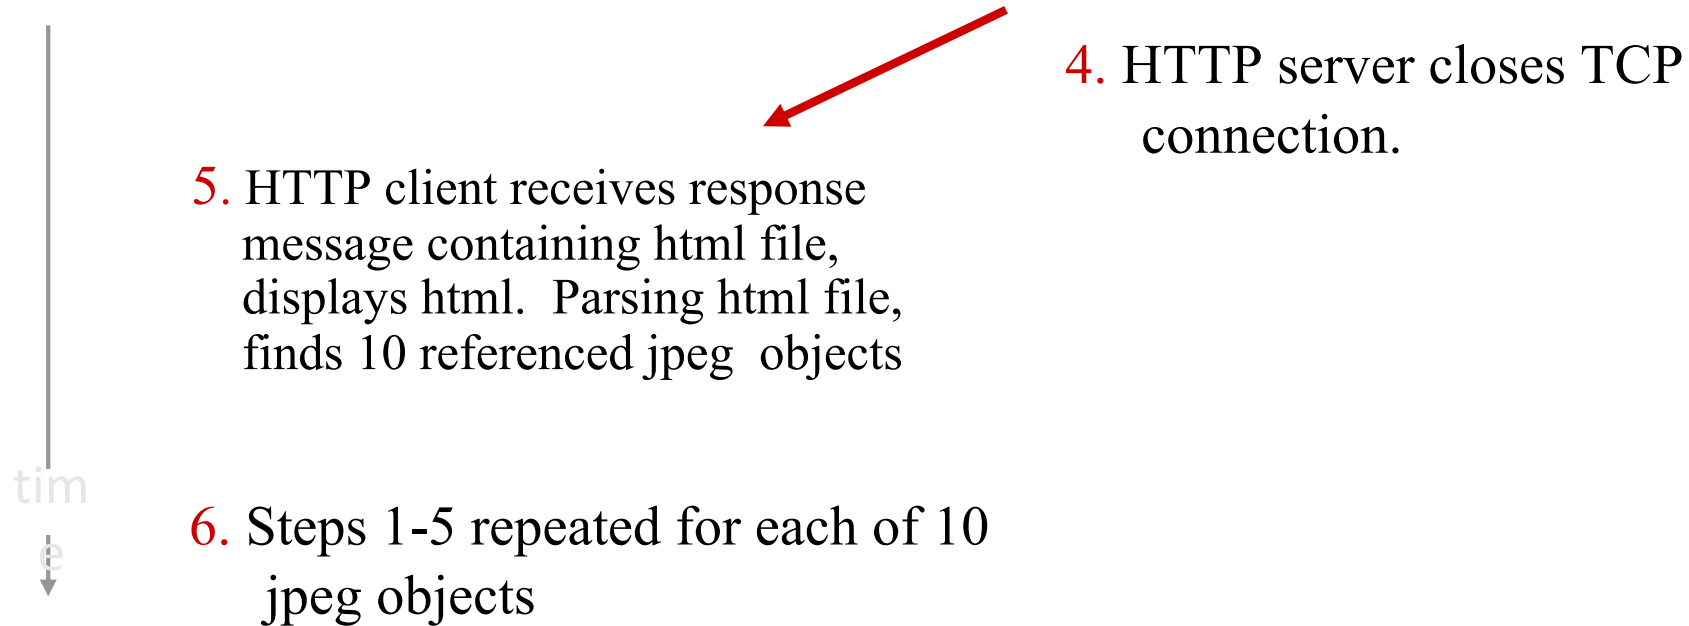
4. HTTP server closes TCP connection.

time  
↓

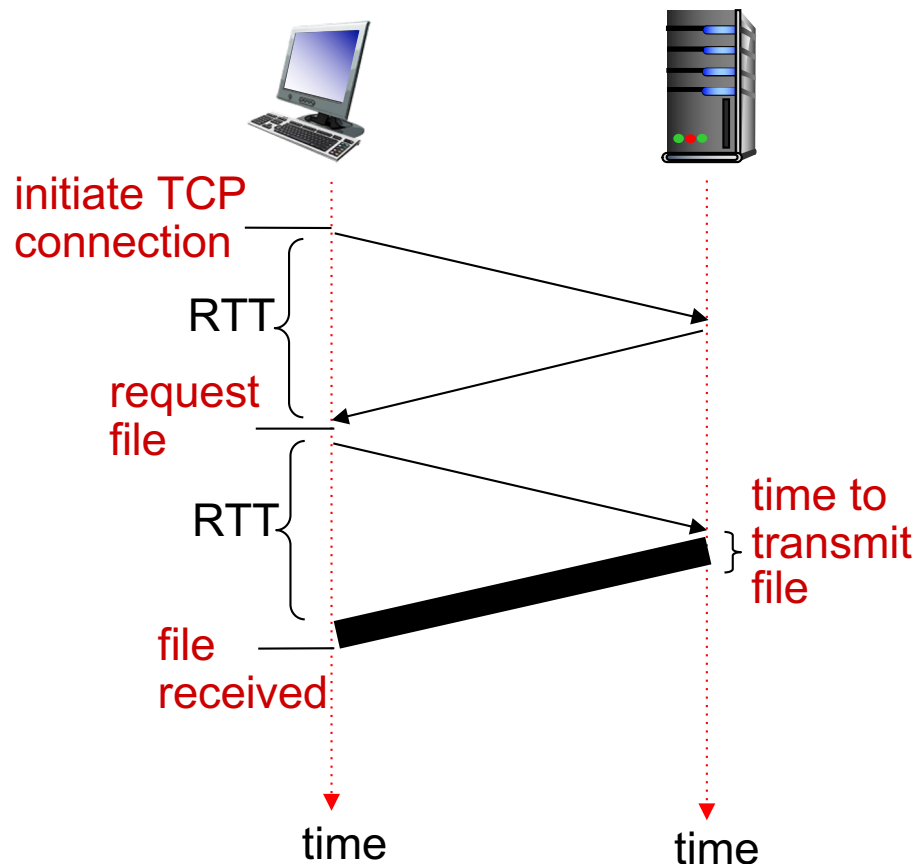
# Non-persistent HTTP



# Non-persistent HTTP



# Non-persistent HTTP: Response Time



**RTT** (definition): **round trip time** for a small packet to travel from client to server and back

## HTTP response time:

- One RTT to initiate TCP connection
- One RTT for HTTP request and first few bytes of HTTP response to return
- File transmission time
- Non-persistent HTTP response time =

$$2\text{RTT} + \text{file transmission time}$$



# Persistent HTTP

## *Non-persistent HTTP issues:*

- Requires 2 RTTs per object
- OS overhead for *each* TCP connection
- Browsers often open parallel TCP connections to fetch referenced objects

## *Persistent HTTP:*

- Server leaves connection open after sending response
- Subsequent HTTP messages between same client/server sent over the open connection
- Client sends requests as soon as it encounters a referenced object
- As little as one RTT for all the referenced objects

# HTTP Request Message

- Two types of HTTP messages: *request, response*
- HTTP request message:
  - ASCII (human-readable format)

request line  
(GET, POST,  
HEAD commands)

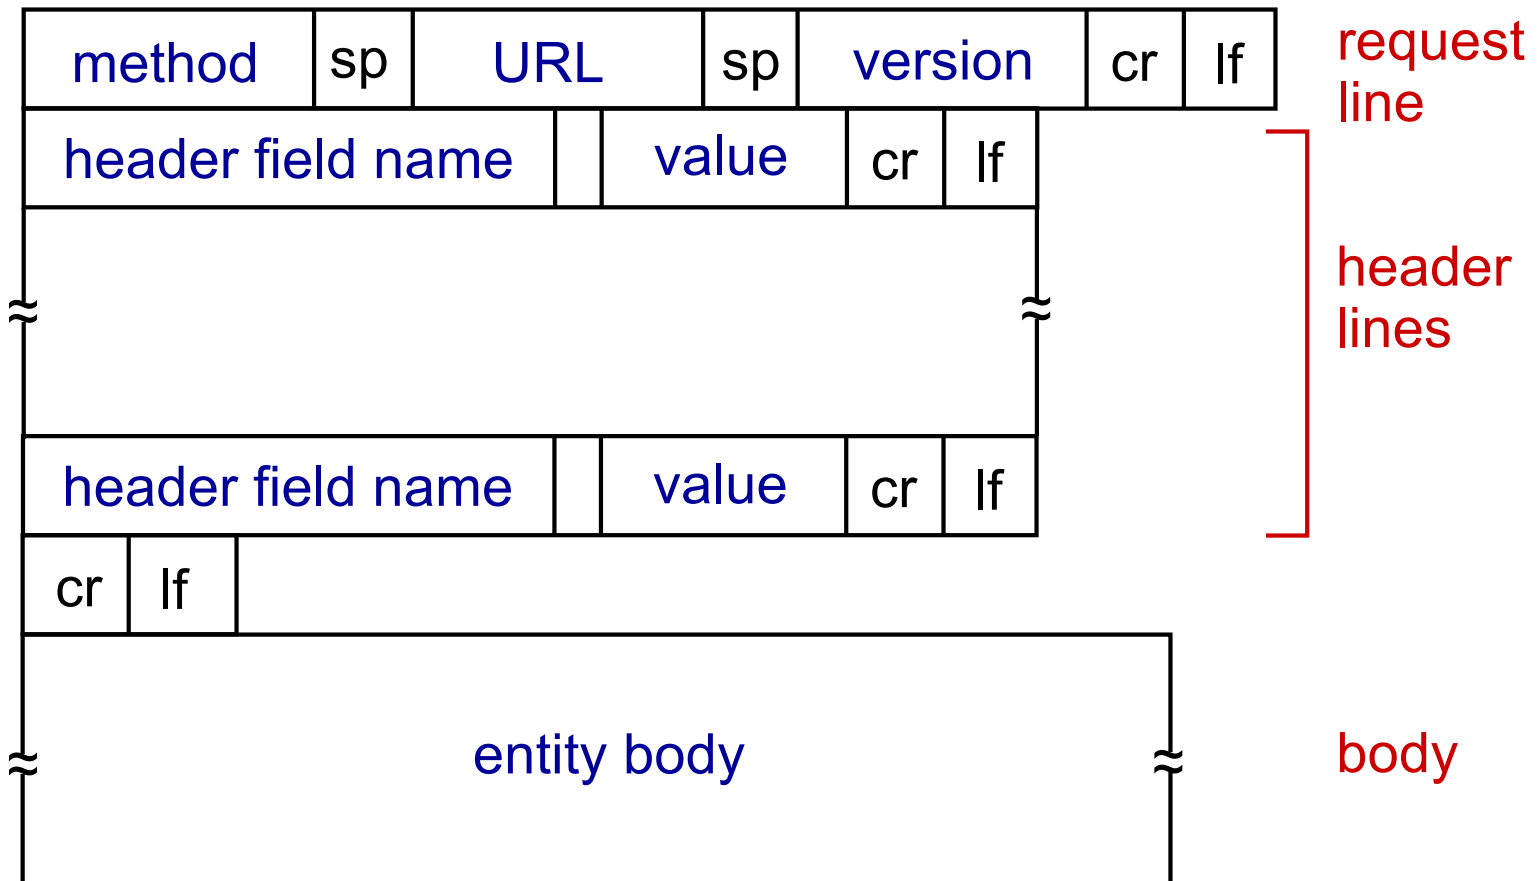
header  
lines

carriage return,  
line feed at start  
of line indicates  
end of header lines

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

carriage return character  
line-feed character

# HTTP Request Message: General Format



# Uploading form Input

## POST method:

- Web page often includes form input
- Input is uploaded to server in entity body

## URL method:

- Uses GET method
- Input is uploaded in URL field of request line:

`www.somesite.com/animalsearch?monkeys&banana`

# Method Types

## HTTP/1.0:

- GET
- POST
- HEAD
  - asks server to leave requested object out of response

## HTTP/1.1:

- GET, POST, HEAD
- PUT
  - uploads file in entity body to path specified in URL field
- DELETE
  - deletes file specified in the URL field

# HTTP Response Message

status line  
(protocol  
status code  
status phrase)

header  
lines

data, e.g.,  
requested  
HTML file

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02
GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html; charset=ISO-8859-1\r\n
\r\n
data data data data data ...
```

# HTTP Response Status Codes

- Status code appears in 1<sup>st</sup> line in server-to-client response message.
- Some sample codes:

## **200 OK**

- request succeeded, requested object later in this msg

## **301 Moved Permanently**

- requested object moved, new location specified later in this msg (Location:)

## **400 Bad Request**

- request msg not understood by server

## **404 Not Found**

- requested document not found on this server

## **505 HTTP Version Not Supported**

# User Server State: cookies

Many Websites use cookies

*four components:*

- 1) Cookie header line of HTTP *response* message
- 2) Cookie header line in next HTTP *request* message
- 3) Cookie file kept on user's host, managed by user's browser
- 4) Back-end database at Web site

**Example:**

- Susan always accesses Internet from her PC
- Visits specific e-commerce site for first time
- When initial HTTP request arrives at site, site creates:
  - Unique ID
  - Entry in backend database for ID

**More detail in tutorial session**



# Cookies: keeping “state” (cont.)


client




server



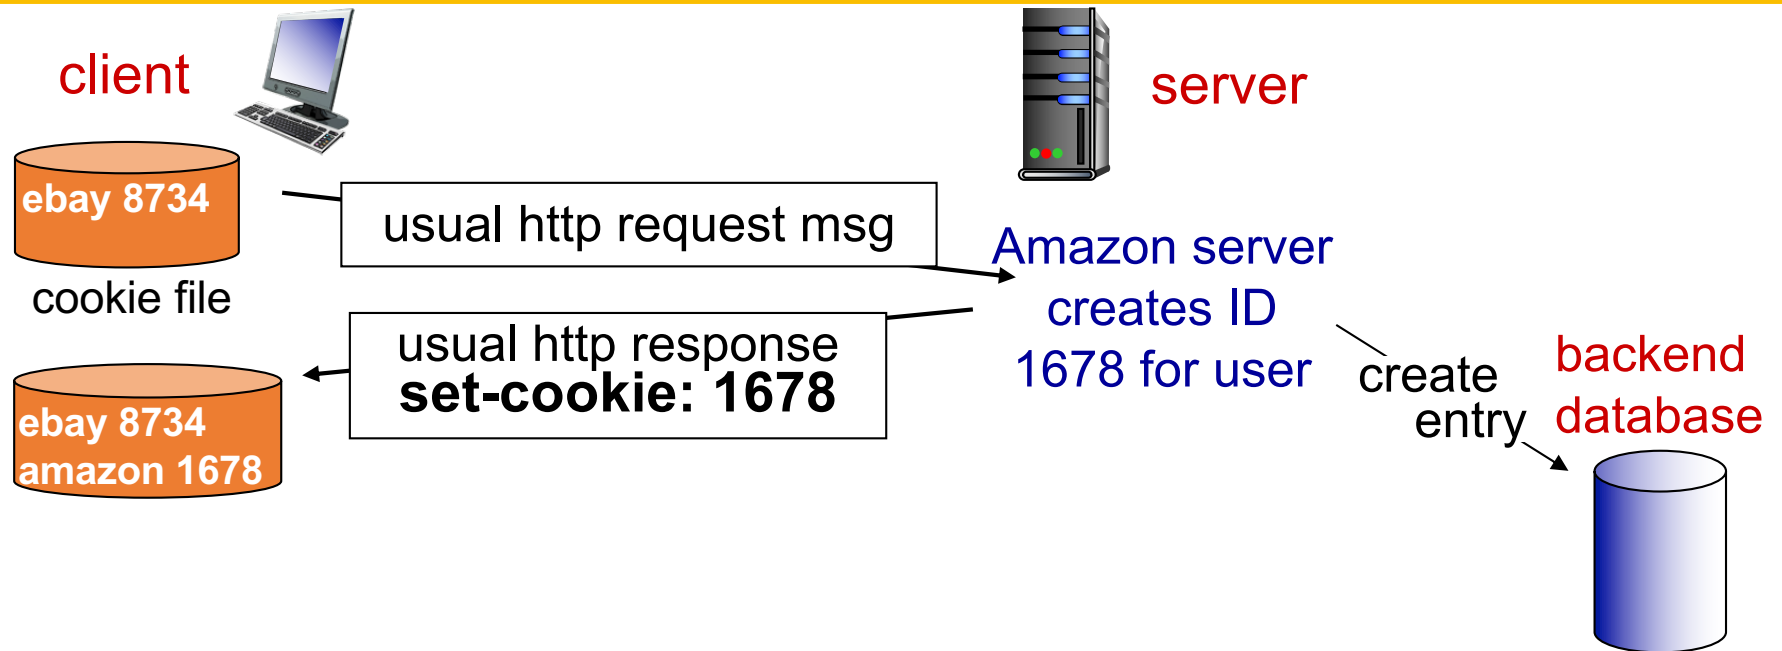
backend  
database



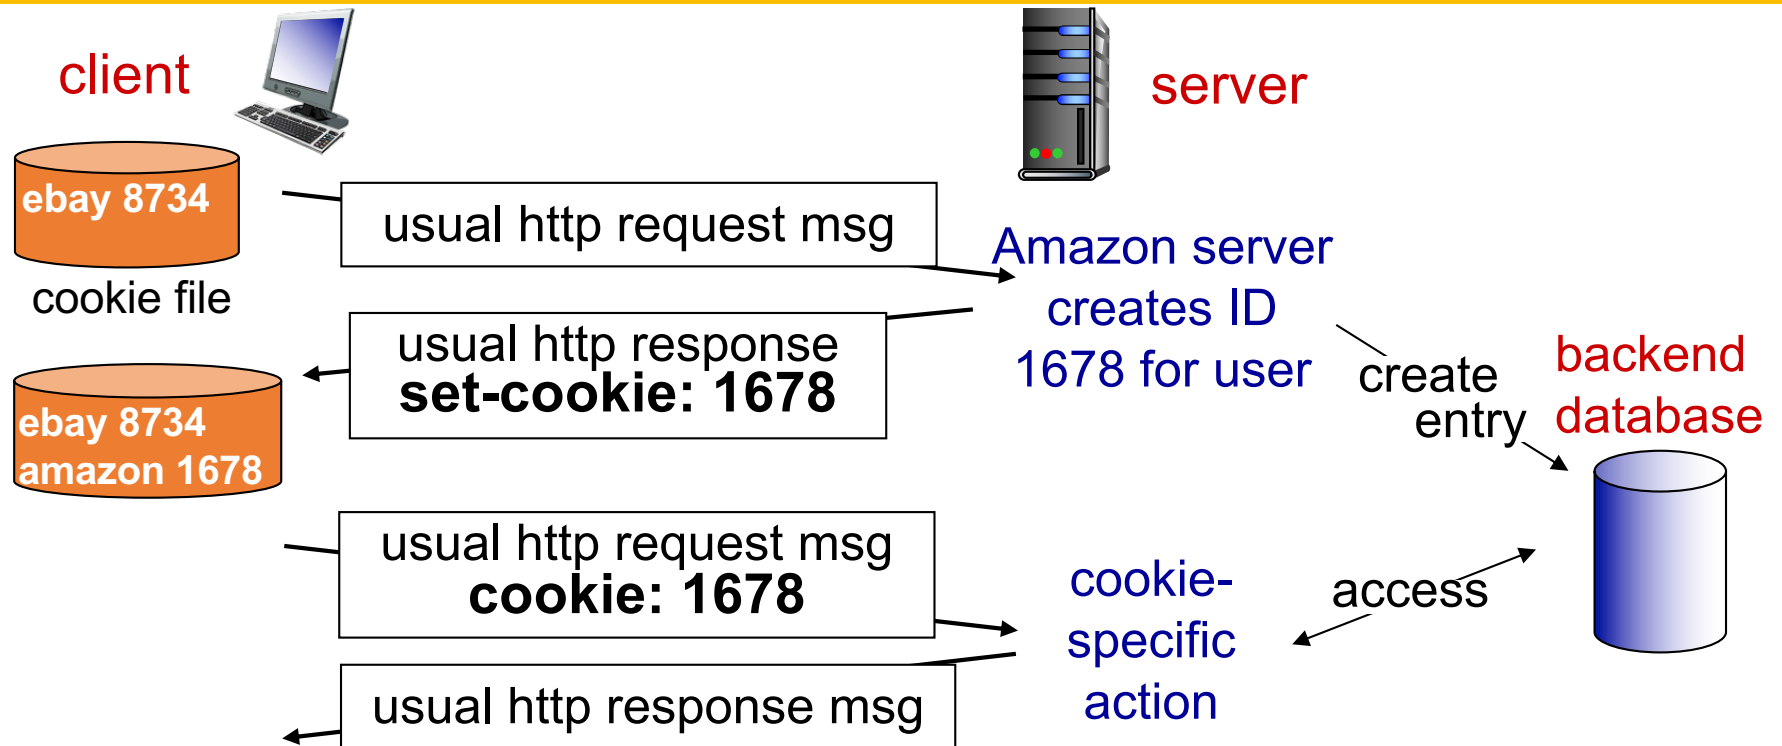
# Cookies: keeping “state” (cont.)



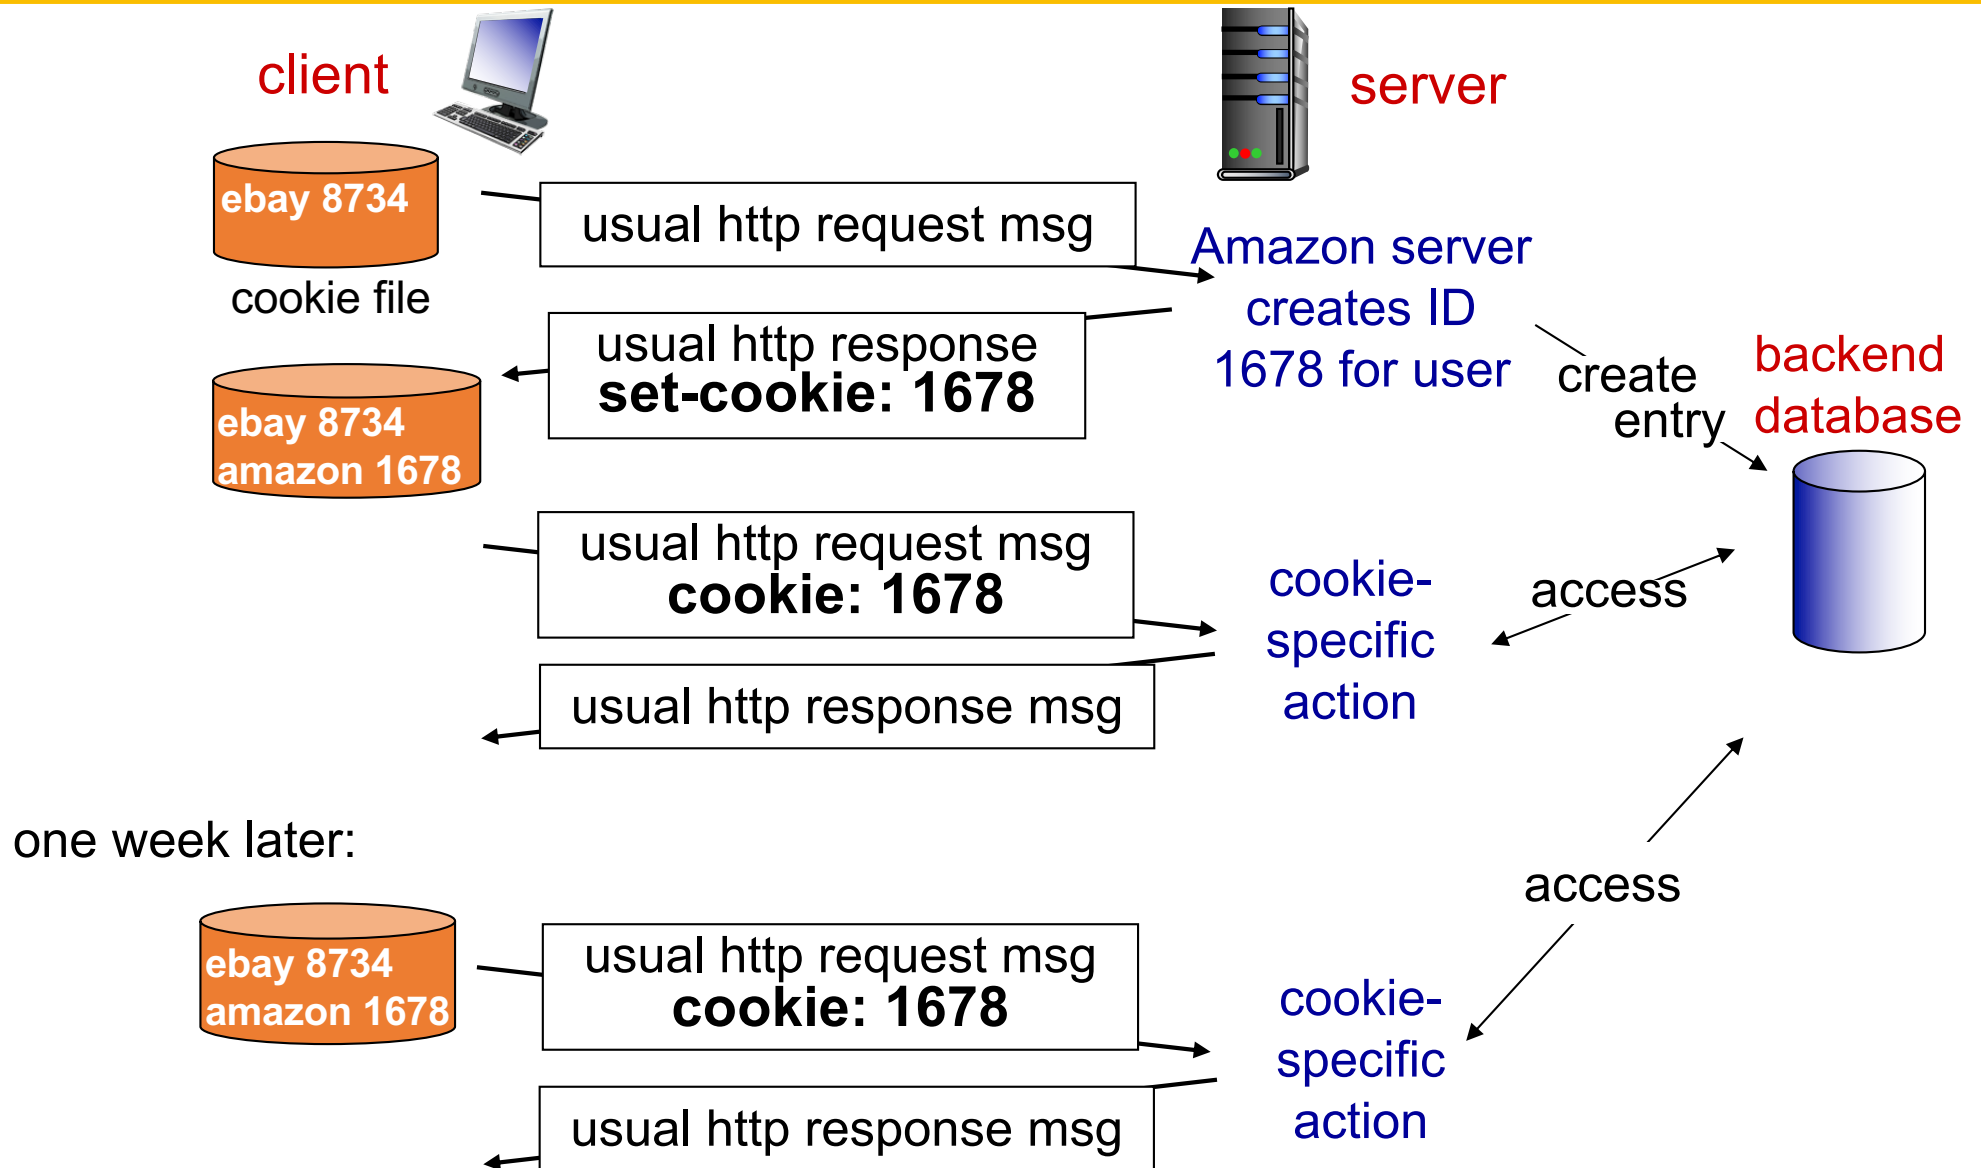
# Cookies: keeping “state” (cont.)



# Cookies: keeping “state” (cont.)



# Cookies: keeping “state” (cont.)



# Cookies (continued)

## *What cookies can be used for:*

- authorization
- Support shopping carts
- recommendations
- user session state (Web e-mail)

## *Cookies and privacy:* aside

- cookies permit sites to learn a lot about you (“user profiling”)
- you may supply name and e-mail to sites


## *How to keep “state”:*

- protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: http messages carry state

# Web cache (or proxy Server)

*Goal:* satisfy client request without involving origin server

- User sets her browser: Web accesses via cache
- Browser sends all HTTP requests to cache
  - If object in cache: then cache returns object
  - else cache requests object from origin server, then returns object to client



# More about Web Caching

- Cache acts as both client and server
  - Server for original requesting client
  - Client to origin server
- Typically the cache is installed by your ISP (eg. university, company, or residential ISP)
- *Why Web caching?*
  - Reduce response time for client requests
  - Reduce traffic on an institution's access link
  - Internet is dense with caches: enables “weak” content providers to effectively deliver content



# Caching Example

## Assumptions:

- avg object size: 1 Mbits
- avg request rate from browsers to origin servers: 15 reqs/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 15 Mbps


## Consequences:

- LAN utilization: 15%
- access link intensity = 99%
- total delay =  
Internet delay + access delay + LAN delay  
= 2 sec + minutes + ~0

Traffic intensity:  $\lambda a / R$

where

$\lambda a$  is the average length and  $R$  the transmission rate




# Caching Example

## Assumptions:

- avg object size: 1M bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 15 Mbps

## Consequences:

- LAN utilization: 15%
- access link intensity = *problem!* 99%
- total delay =  
Internet delay + access delay + LAN delay  
= 2 sec + minutes + ~0




# Caching Example: fatter access link

## Assumptions:

- avg object size: 1M bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 15 Mbps

## Consequences:

- LAN utilization: 15%
- access link intensity = 99%
- total delay =  
Internet delay + access delay + LAN delay  
= 2 sec + minutes + ~0



# Caching Example: fatter access link


## Assumptions:

- avg object size: 1M bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: ~~15 Mbps~~ → 100 Mbps

## Consequences:

- LAN utilization: 15%
- access link intensity = ~~99%~~ → 9.9%
- total delay = Internet delay + access delay + LAN delay  

$$= 2 \text{ sec} + \text{minutes} \rightarrow \text{msecs}$$



*Cost:* increased access link speed (not cheap!)


# Caching example: install local cache

## Assumptions

- Avg object size: 1M bits
- Avg request rate from browsers to origin servers: 15/sec
- Avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- Access link rate: 15 Mbps

## Consequences

- LAN utilization: 15%
- Access link intensity =
- Total delay =




# Caching example: install local cache

## Assumptions

- Avg object size: 1M bits
- Avg request rate from browsers to origin servers: 15/sec
- Avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- Access link rate: 15 Mbps

## Consequences

- LAN utilization: 15%
- Access link intensity =
- Total delay =



*Cost:* web cache (cheap!)

# Caching example: install local cache

## Assumptions


- Avg object size: 1M bits
- Avg request rate from browsers to origin servers: 15/sec
- Avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- Access link rate: 15 Mbps

## Consequences

- LAN utilization: 15%
- Access link intensity = ?
- Total delay = ?

*How to compute link intensity, delay?*


*Cost:* web cache (cheap!)



# Caching example: install local cache

*Calculating access link intensity,  
delay with cache:*

- Suppose cache hit rate is 0.4
  - 40% requests satisfied at cache,  
60% requests satisfied at origin






# Caching example: install local cache

## *Calculating access link intensity, delay with cache:*


- Suppose cache hit rate is 0.4
  - 40% requests satisfied at cache, 60% requests satisfied at origin
- Access link intensity:
  - 60% of requests use access link
- Data rate to browsers over access link =  $0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$ 
  - intensity =  $0.9 / 1.54 = .58$



# Caching example: install local cache

*Calculating access link intensity,  
delay with cache:*

- Suppose cache hit rate is 0.4
  - 40% requests satisfied at cache, 60% requests satisfied at origin
- Access link intensity:
  - 60% of requests use access link
- Data rate to browsers over access link =  $0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$ 
  - intensity =  $0.9 / 1.54 = .58$
- total delay
  - =  $0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
  - =  $0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$
  - less than with 100 Mbps link (and cheaper too!)

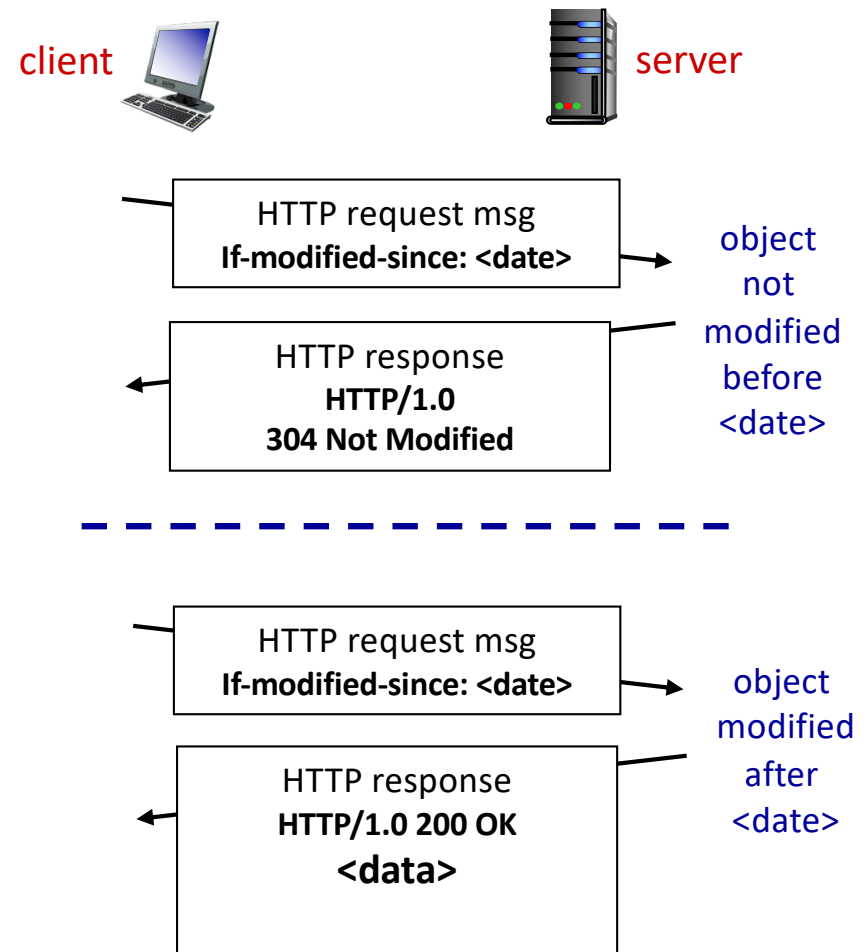


See 2.2.5 in Kurose 8ed

# Conditional GET

**Goal:** don't send object if cache has up-to-date cached version

- no object transmission delay (or use of network resources)
- **client:** specify date of cached copy in HTTP request  
**If-modified-since: <date>**
- **server:** response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not Modified**



# HTTP/2

*Key goal:* decreased delay in multi-object HTTP requests

HTTP1.1: introduced **multiple, pipelined GETs** over single TCP connection

- server responds *in-order* (FCFS: first-come-first-served scheduling) to GET requests
- with FCFS, small object may have to wait for transmission (**head-of-line (HOL) blocking**) behind large object(s)
- loss recovery (retransmitting lost TCP segments) stalls object transmission

# HTTP/2

*Key goal:* decreased delay in multi-object HTTP requests


HTTP/2: [RFC 7540, 2015] increased flexibility at *server* in sending objects to client:

- methods, status codes, most header fields unchanged from HTTP 1.1
- transmission order of requested objects based on client-specified object priority (not necessarily FCFS)
- *push* unrequested objects to client
- divide objects into frames, schedule frames to mitigate HOL blocking


# The Application Layer

- Principles of network applications
- Web and HTTP
- FTP

# File Transfer Protocol (FTP)



# FTP: separate control, data connections



- FTP client contacts FTP server at port 21, using TCP
- Client authorized over control connection
- Client browses remote directory, sends commands over control connection
- When server receives file transfer command, **server** opens 2<sup>nd</sup> TCP data connection (for file) to client
- After transferring one file, server closes data connection
- FTP server maintains “state”: current directory, earlier authentication



# Summary

- Computer Networks
- Types of Networks
- Open Systems Interconnection (OSI) Model
- The Application Layer
- Application Architecture
- Principles of Network Applications
- Web and HTTP
- FTP